



Microsoft®
Research

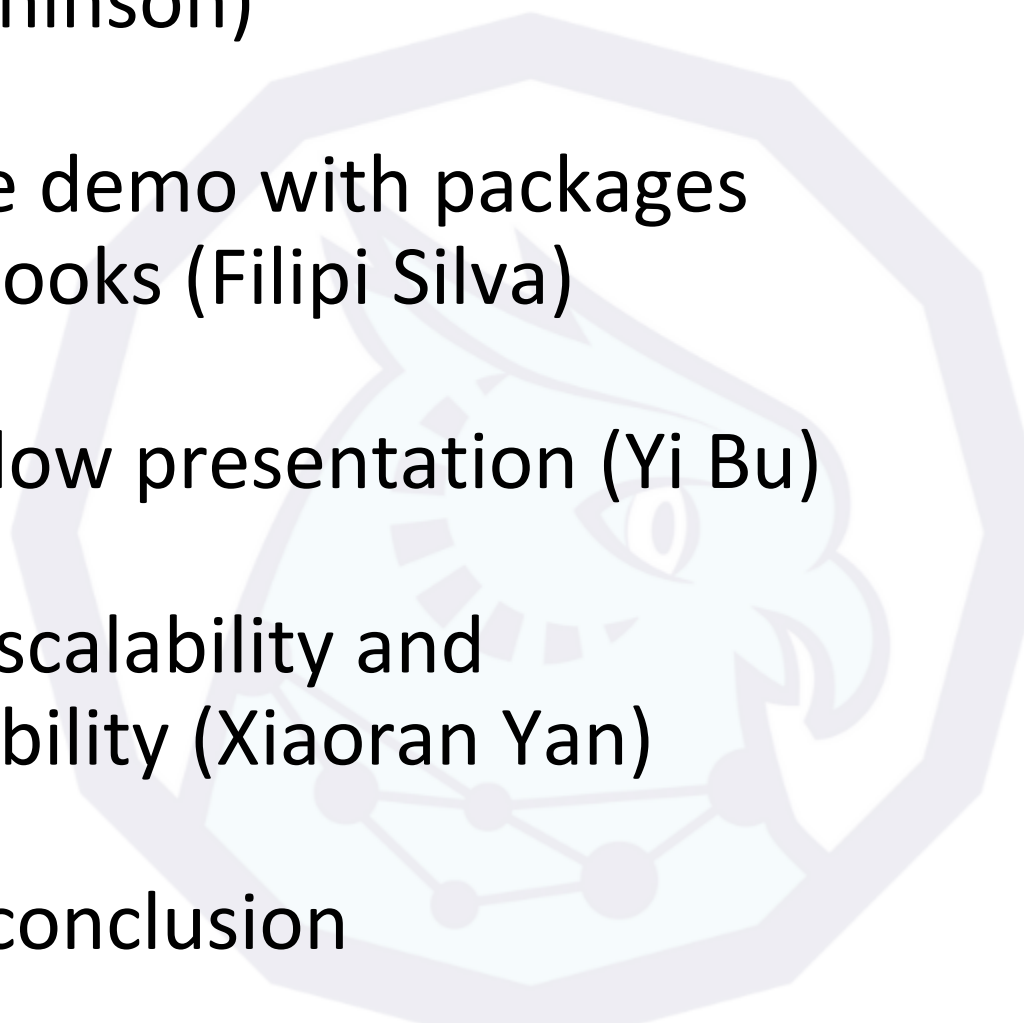


Hands-on Tutorial

Supported by Microsoft Research



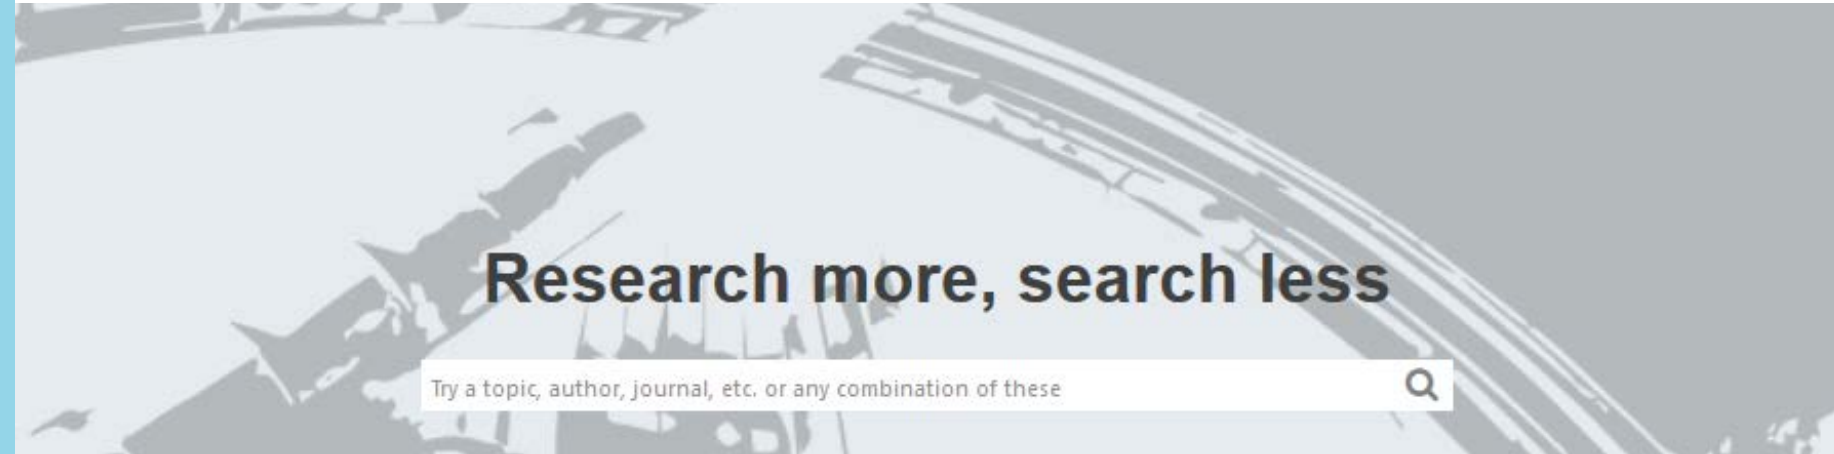
Program overview

- The CADRE project (Val Pentchev)
 - Hands on intro to CADRE (Mat Hutchinson)
 - Interactive demo with packages and notebooks (Filipi Silva)
 - CADRE fellow presentation (Yi Bu)
 - Demo for scalability and Reproducibility (Xiaoran Yan)
 - Q&A and conclusion
- 

<https://academic.microsoft.com/home>

Microsoft Academic Graph

An Overview of Microsoft Academic Service (MAS) and Applications. In *Proceedings of the 24th International Conference on World Wide Web 2015*.



Publications

210,365,701

Coming soon



Authors

254,317,172

Learn more



Fields of Study

229,763

Learn more



Conferences

4,341

Learn more



Journals

48,659

Learn more



Institutions

25,439

Learn more



Microsoft®
Research




The CADRE project


Val Pentchev



The CADRE team



CADRE Leadership



Partners



University of Iowa Libraries



University of Michigan Libraries



Michigan State University Libraries



University of Minnesota Libraries



Ohio State University Libraries



Penn State University Libraries



Purdue University Libraries



Rutgers University Libraries



Health Partners



Pervasive Technology Institute



Midwest Big Data Hub



South Big Data Hub



West Big Data Hub



Microsoft Research



Web of Science Group

This project was made possible in part by the Institute of Museum and Library Services LG-70-18-0202.

Topic 1

- Content



Topic 2

Content

- Content





Microsoft®
Research



Hands on intro to CADRE

Mat Hutchinson



Demo 1

<https://github.com/iuni-cadre/ISSI-tutorial>



Questions?





Microsoft®
Research



Interactive demo

Filipi Silva



Demo 2

<https://github.com/iuni-cadre/ISSI-tutorial>



Demo 3

<https://github.com/iuni-cadre/ISSI-tutorial>



Questions?





Microsoft®
Research



CADRE Fellows

Xiaoran Yan




CADRE related events

Apr. 2019



- 2019 CADRE meeting
- CADRE Fellowship open
- 1st Fellows announced
- ISSI workshop & tutorial



Sep. 2019



SAPIENZA
UNIVERSITÀ DI ROMA

May. 2020



- 2020 CADRE meeting
- BTAA Library Conference 2020
- 2020 CADRE hack-a-thon



INDIANA UNIVERSITY
BLOOMINGTON

CADRE Fellowship program

- Gain access to the big bibliometric data sets
- Receive data and technical support for your project
- Join the CADRE community on Slack channels, GitHub repositories and other platforms
- Have early access to free cloud computing resources
- Receive travel scholarships

Utilizing Data Citation for Aggregating, Contextualizing, and Engaging with Research Data in STEM Education Research

Researchers: Michael Witt, Loran Carleton Parker, Ann Bessenbacher

Affiliation: Purdue University



MCAP: Mapping Collaborations and Partnerships in SDG Research

Researchers: Jane Payumo, Devin Higgins, Scout Calvert, Guangming He
Affiliation: Michigan State University



The global network of air links and scientific collaboration – a quasi-experimental analysis

Researchers: Katy Börner, Adam Ploszaj, Lisel Record, Bruce Herr II

Affiliation: Indiana University Bloomington and University of Warsaw



Measuring and Modeling the Dynamics of Science Using the CADRE Platform

Researchers: Russell Funk, Michael Park, Thomas Gebhart, Britta Glennon, Julia Lane, Raviv Murciano-Goroff, Matthew Ross, Jina Lee, Erin Leahey

Affiliation: University of Minnesota, University of Pennsylvania, New York University, Boston University, University of Arizona



Comparative analysis of legacy and emerging journals in mathematical biology

Researchers: Marisa Conte, Samuel Hansen, Scott Martin, Santiago Schnell

Affiliation: University of Michigan and University of Michigan Medical School



Systematic over-time study of the similarities and differences in research across mathematics and the sciences

Researcher: Samuel Hansen
Affiliation: University of Michigan



A user story from CADRE fellows



Understanding citation impact of scientific publications through ego-centered citation networks

Researchers: Yi Bu, Chao Min, Ying Ding

Affiliation: Indiana University Bloomington and Nanjing University





Microsoft®
Research



Exploring ego-centered citation networks: A technical introduction

Yi Bu¹, Chao Min², and Ying Ding¹

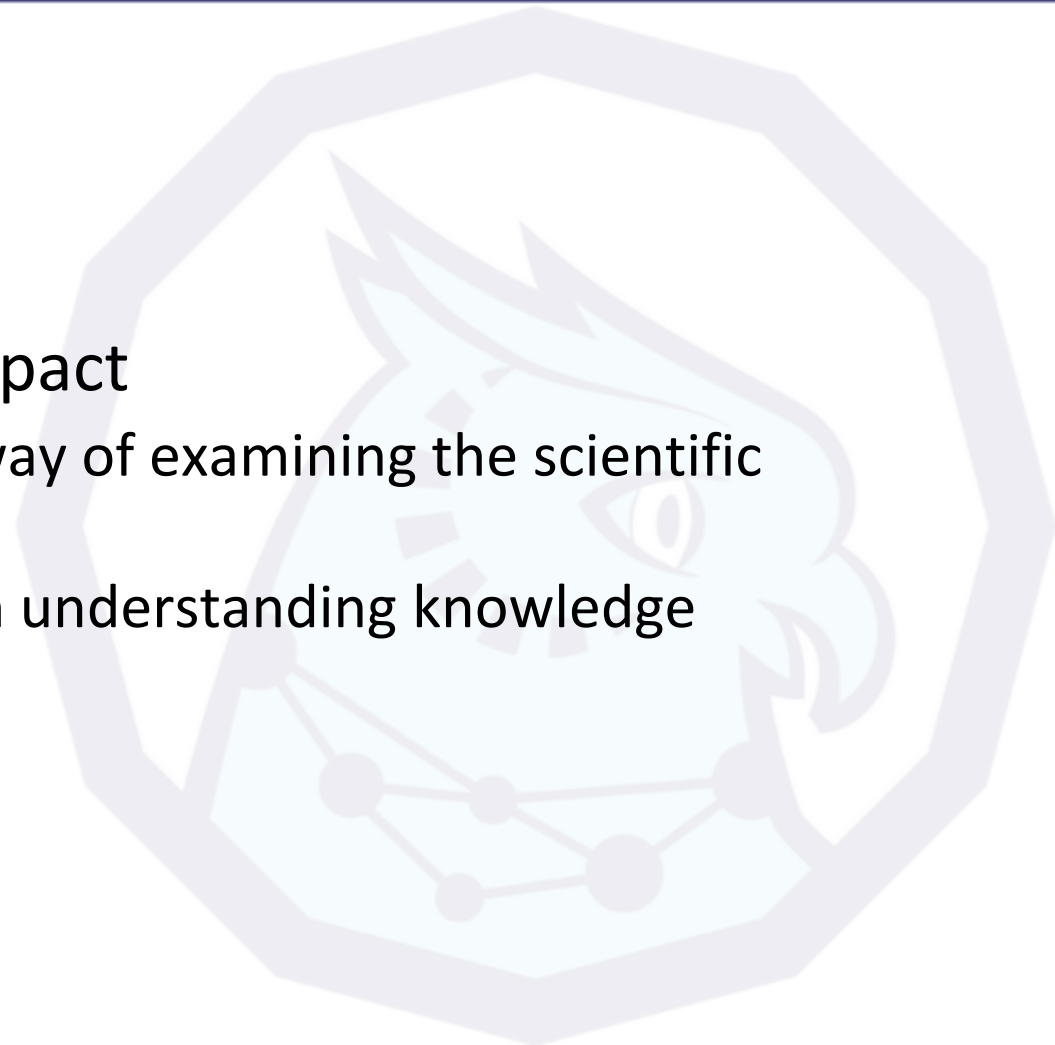
1: School of Informatics, Computing, and Engineering, Indiana University, U.S.A.

2: School of Information Management, Nanjing University, China




Understanding citation impact of scientific publications

- Citation impact as a type of impact
 - ✓ Citation impact among all types of impact
 - ✓ Citation impact of scientific publications
- Benefits from understanding citation impact
 - ✓ Measuring citation impact offers a useful way of examining the scientific impact of a publication.
 - ✓ Measuring citation impact can also assist in understanding knowledge diffusion and the use of information.





Understanding citation impact of scientific publications (cont.)

- Previous ways of understanding citation impact of scientific publications:
 - ✓ Count-based strategies: raw citation count, normalized citation measures...
 - ✓ Network-based strategies: PageRank, EigenFactor...



Understanding citation impact of scientific publications (cont.)

- Local details are missing!
 - ✓ “Deep” or “wide” impact?




Understanding citation impact of scientific publications (cont.)

- Local details are missing!
 - ✓ How does an article impact other research, and what are the patterns? The direct citations between citing publications (DCCPs) offer a good way to mine how a publication impacts other research.


Cited publication	citing publication						
		SSH	BHS	PSE	LES	MCS	subtotal
	SSH	11138	224	16	5	37	11420
	BHS	440	1254	2	11	1	1708
	PSE	137	1	19	3	18	178
	LES	57	13	3	11	0	84
	MCS	194	0	17	0	26	237
	subtotal	11966	1492	57	30	82	13627

year	SSH	BHS	PSE	LES	MCS
2006	13	0	0	0	0
2007	111	0	0	0	0
2008	455	0	2	2	4
2009	753	9	3	0	0
2010	1155	19	0	1	0
2011	1310	80	2	1	12
2012	1092	39	3	1	9
2013	1440	187	19	3	41
2014	1110	449	30	2	31
2015	1161	361	12	12	13
2016	1491	290	44	57	60
2017	1329	274	63	5	67


Understanding citation impact of scientific publications (cont.)



Ego-centered citation networks as a tool to understand citation impact



(a)



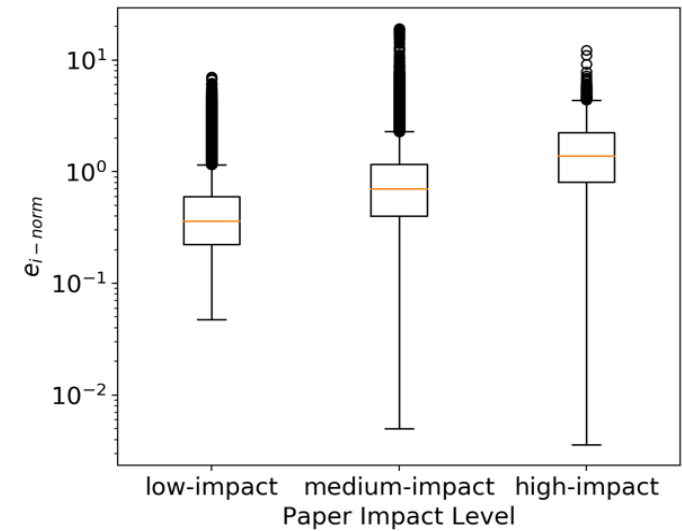


(b)

Preliminary research questions






- Do DCCPs occur frequently?
- How does DCCPs different in papers with different citation impacts and in different years?



Preliminary results: The universality of DCCPs



Preliminary results (cont.)



Technical details: Extracting citing relationships from the raw WoS tables

- SQL extraction as a .txt file:

```
import psycopg2
conn = psycopg2.connect(database = 'core_data', user = 'buyi', password = )
cur = conn.cursor()
cur.execute("SELECT paper_id, paper_reference_id FROM mag_core.paper_references;")
outFile = open("mag_citing.txt", "w+")
lines = ['citing id=====cited id']
for row in cur:
    if str(row[0]) in paper_id_set and str(row[1]) in paper_id_set:
        lines.append('{:}====={:}'.format(str(row[0]), str(row[1])))
    if len(lines) % 100000 == 0:
        outFile.write('\n'.join(lines) + '\n')
        lines = []

outFile.write('\n'.join(lines) + '\n')
cur.close()
```

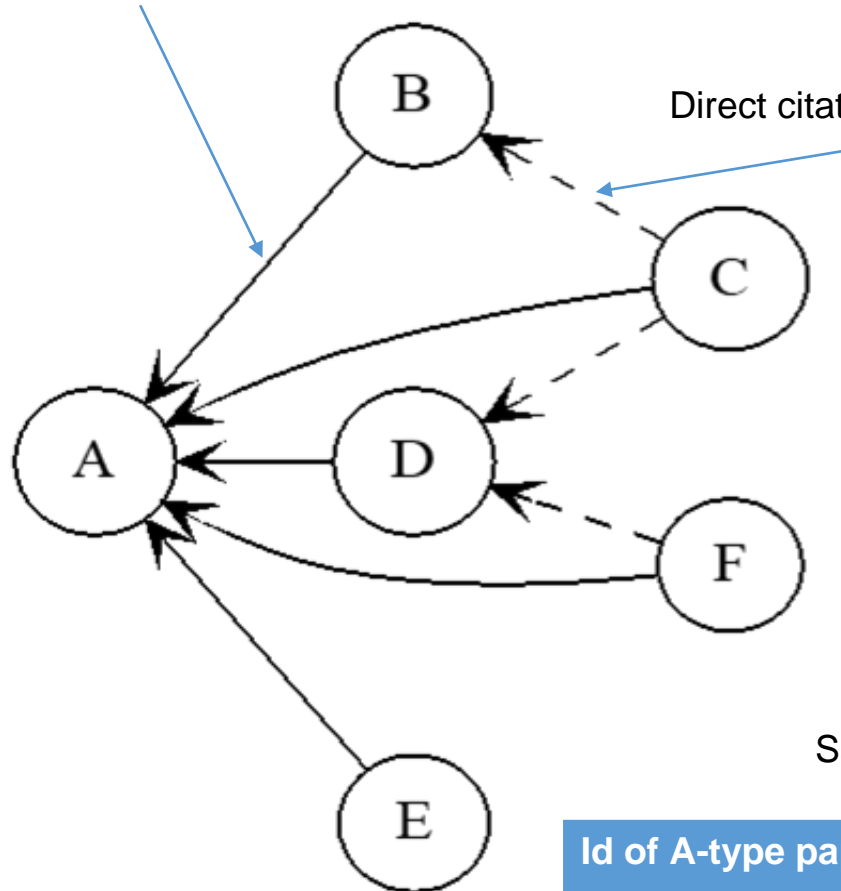
- .txt file to a Python dictionary:

✓ If paper in paper_citing.keys()



Difficulty 1: How to extract DCCPs?

Direct citations to A



Direct citations between citing publications (from the perspective of A)

Sample output:

Id of A-type paper (focal)	Id of B-type paper	Id of C-type paper

Difficulty 1: How to extract DCCPs? (cont.)

- This task is computationally expensive:
 - ✓ In MAG, we have ~0.1 billion papers. The below Python script will perhaps take forever...

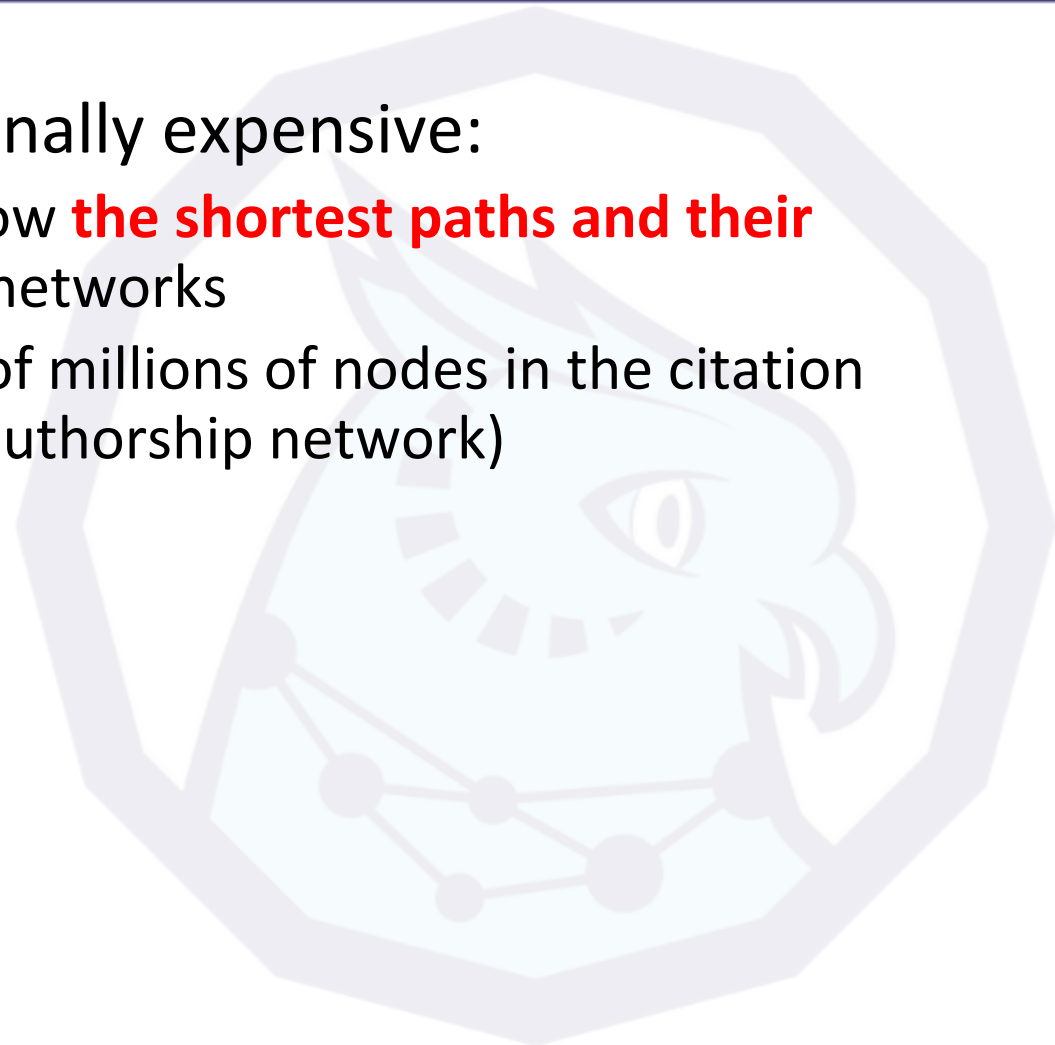
```
indirect_citation = defaultdict(list)
for paper in paper_year.keys(): # for papers that have pub_year information
    for citing_paper_1 in paper_citing[paper]:
        for citing_paper_2 in paper_citing[paper]:
            if citing_paper_1 in paper_citing[citing_paper_2]:
                temp = []
                temp.append(citing_paper_1)
                temp.append(citing_paper_2)
                indirect_citation[paper].append(temp)
```


Difficulty 2: Self-citations in ego-centered citation networks?

- If two papers (A and B) share at least one co-author and B cites A, such citation is called a self-citation (first-order self-citation).
- How about these circumstances, when B cites A?
 - ✓ A and B don't share co-authors, but A and C do, and B and C do. [second-order self-citations]
 - ✓ A and B don't share co-authors, but A and C do, B and D do, and C and D do. [third-order self-citations]
 - ✓ This indicates how researchers' social distance impacts on their self-citation patterns.
- **How to technically achieve these?**

Difficulty 2: Self-citations in ego-centered citation networks?

- Completing this task is also computationally expensive:
 - ✓ Deriving n-order self-citations need to know **the shortest paths and their lengths** in the co-authorship and citation networks
 - ✓ Such networks are quite huge (hundreds of millions of nodes in the citation network, and millions of nodes in the co-authorship network)



Questions?

Presenter: Yi Bu, Indiana University

Email: buyi@iu.edu

Website: <https://buyi08.wixsite.com/yi-bu>





Microsoft®
Research



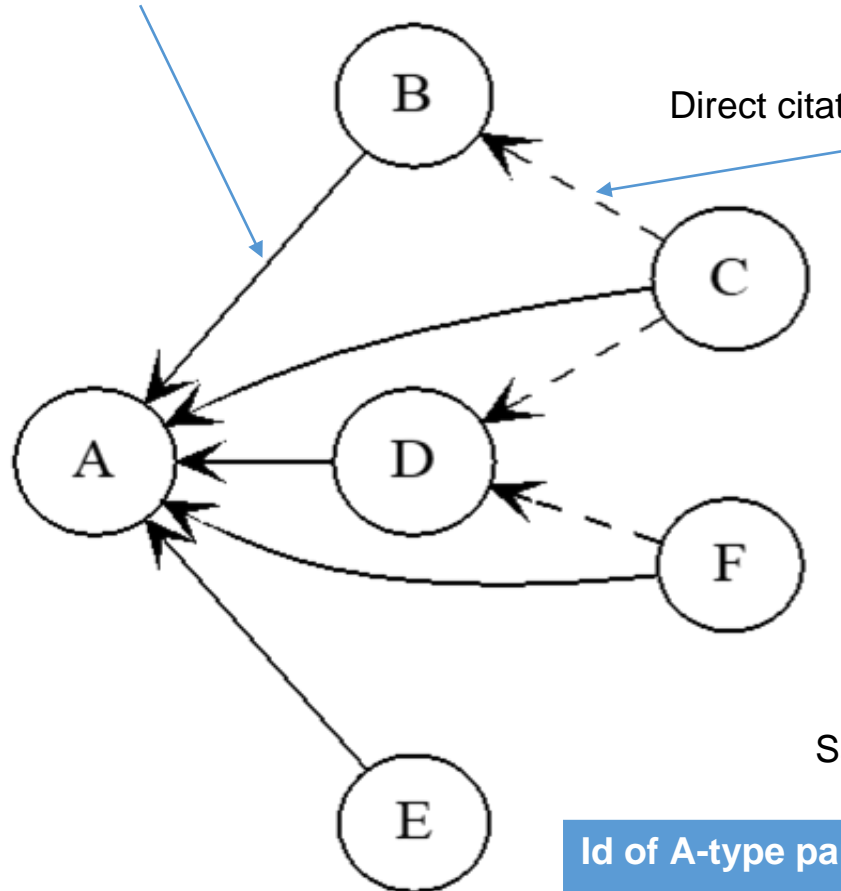
Scalability & Reproducibility

Xiaoran Yan



Difficulty 1: How to extract DCCPs?

Direct citations to A



Direct citations between citing publications (from the perspective of A)

Sample output:

Id of A-type paper (focal)	Id of B-type paper	Id of C-type paper

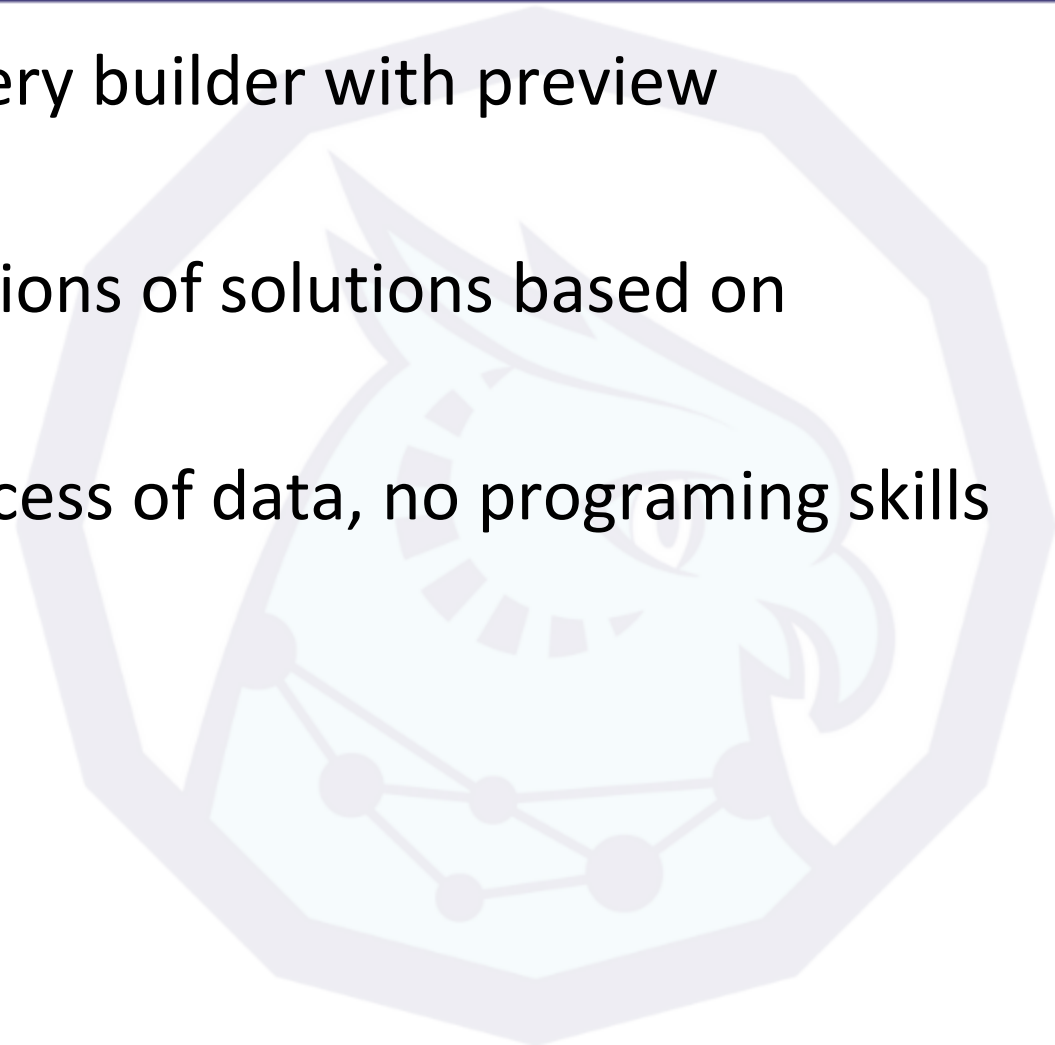
Difficulty 1: How to extract DCCPs? (cont.)

- This task is computationally expensive:
 - ✓ In MAG, we have ~0.1 billion papers. The below Python script will perhaps take forever...

```
indirect_citation = defaultdict(list)
for paper in paper_year.keys(): # for papers that have pub_year information
    for citing_paper_1 in paper_citing[paper]:
        for citing_paper_2 in paper_citing[paper]:
            if citing_paper_1 in paper_citing[citing_paper_2]:
                temp = []
                temp.append(citing_paper_1)
                temp.append(citing_paper_2)
                indirect_citation[paper].append(temp)
```

CADRE's solution

- An easy to use graphical interface of a query builder with preview functionality
- A unified engine with optimized combinations of solutions based on relational/graph/document databases
- For users who want intuitive and quick access of data, no programming skills required
- In development: APIs for power users



CADRE's solution



Access over 220 million scientific publications



Effortlessly query data and analyze results



Reproduce research & leverage tools

CADRE's solution



Databases

GUI-query

Notebooks

RAC

Demo 4

<https://github.com/iuni-cadre/ISSI-tutorial>



Questions?

Presenter: Xiaoran Yan, Indiana University

Email: yan30@iu.edu



CADRE's solution



Access over 220 million scientific publications

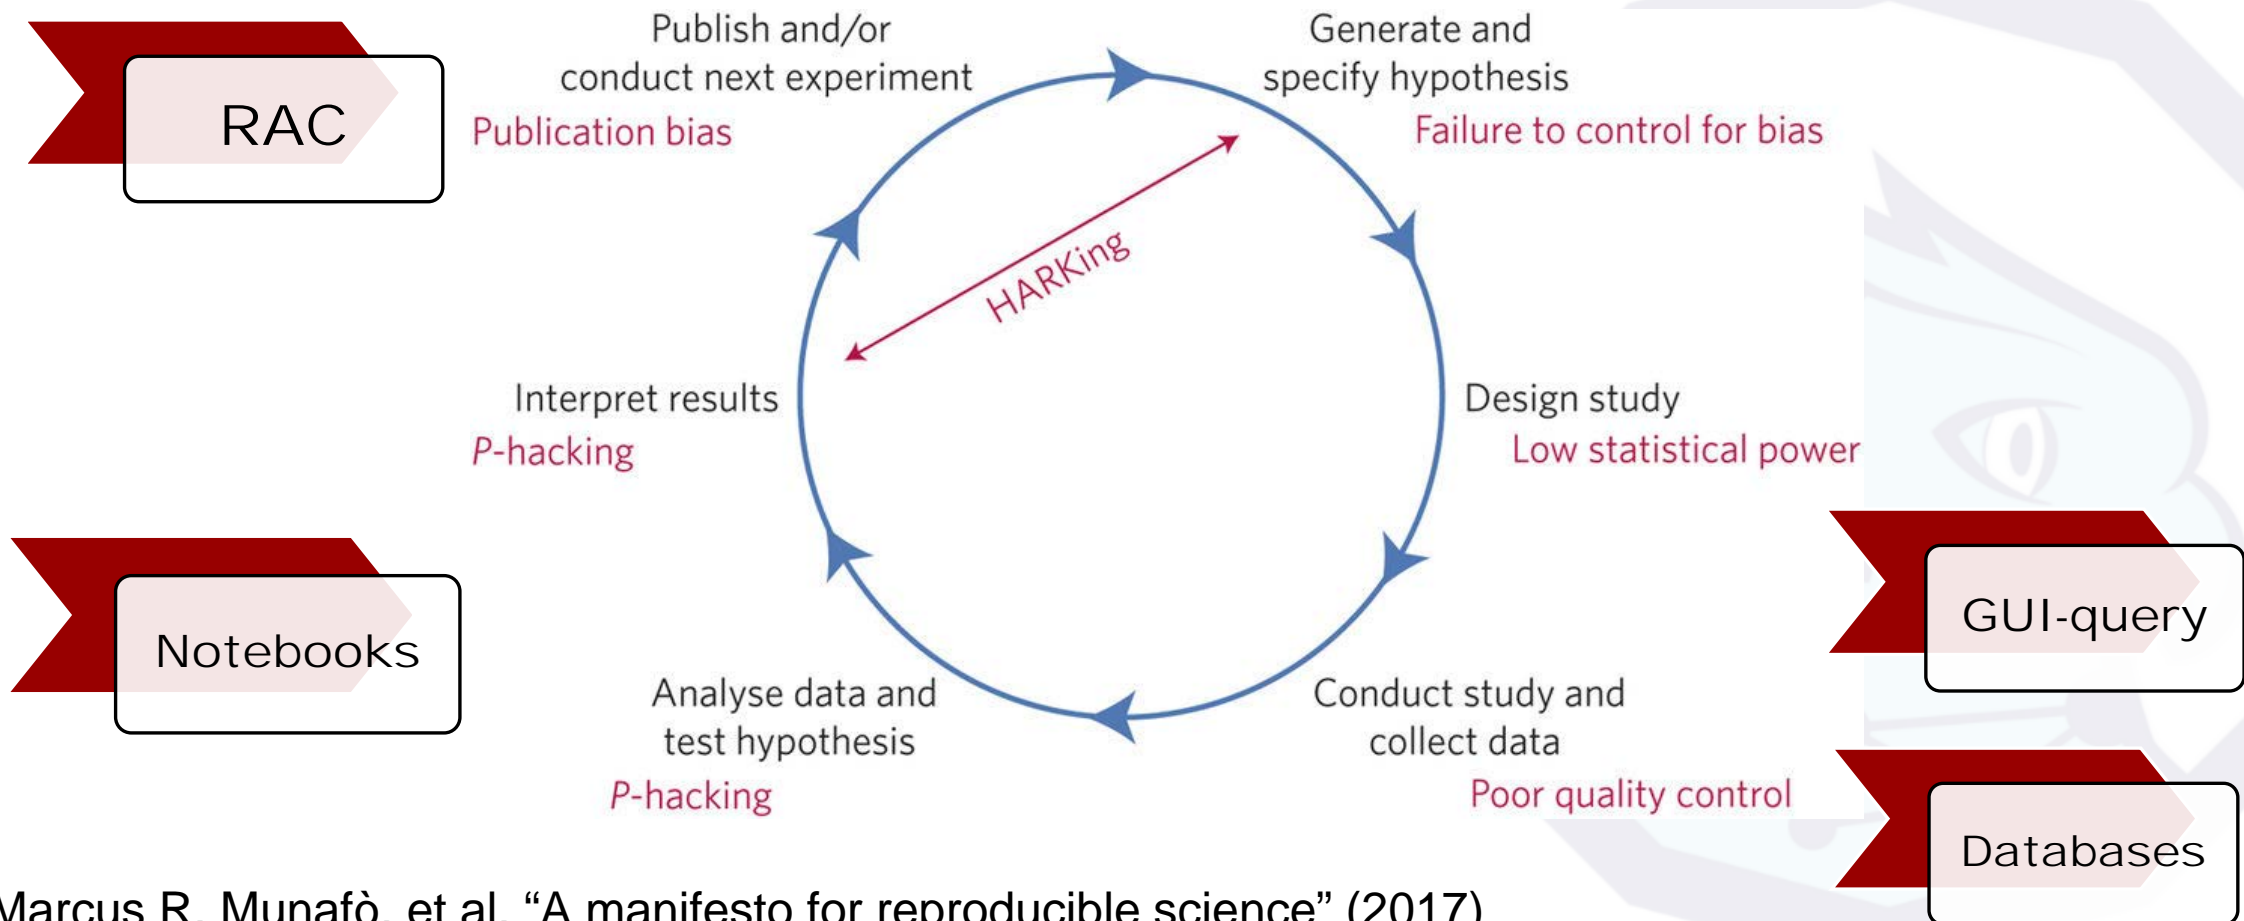


Effortlessly query data and analyze results



Reproduce research & leverage tools

The reproducibility “Crisis”



Marcus R. Munafò, et al. “A manifesto for reproducible science” (2017)

Spectrum of Reproducibility



Computational



Statistical



Empirical



Stodden, Victoria. "Resolving Irreproducibility in Empirical and Computational Research" (2013)

Current solutions

IU-AMBITION / MASS

<> Code

Issues 0

Pull requests 0

Projects 0

Matlab code of minimum absolute spectral similarity

6 commits

1 branch

Branch: master

New pull request

everyxs Update README.md	
LFR.mat	Add files via upload
README.md	Update README.md
absSpecSim.m	Add files via upload
main.m	Add files via upload
sparsify.py	Added a Python implementation

README.md


Matlab function for the Minimum absolute spectral similarity (MASS)

Publications

15. Yan, X., Jeub, L., Flammini, A., Radicchi, F., and Fortunato, S. Weight Thresholding on Complex Networks. *Accepted by Physical Review E, Issue/Batch: 10_05_18 (2018)*
<https://arxiv.org/abs/1806.07479>
Source code: <https://github.com/IU-AMBITION/MASS>
14. Faskowitz, J., Yan, X., Zuo, X.-N., and Sporns, O. Scientific reports, 8(1):12997
<https://www.nature.com/articles/s41598-018-31202-1>



Big data pipelines in the industry



TensorFlow + kubernetes



Kubeflow

CADRE's solution



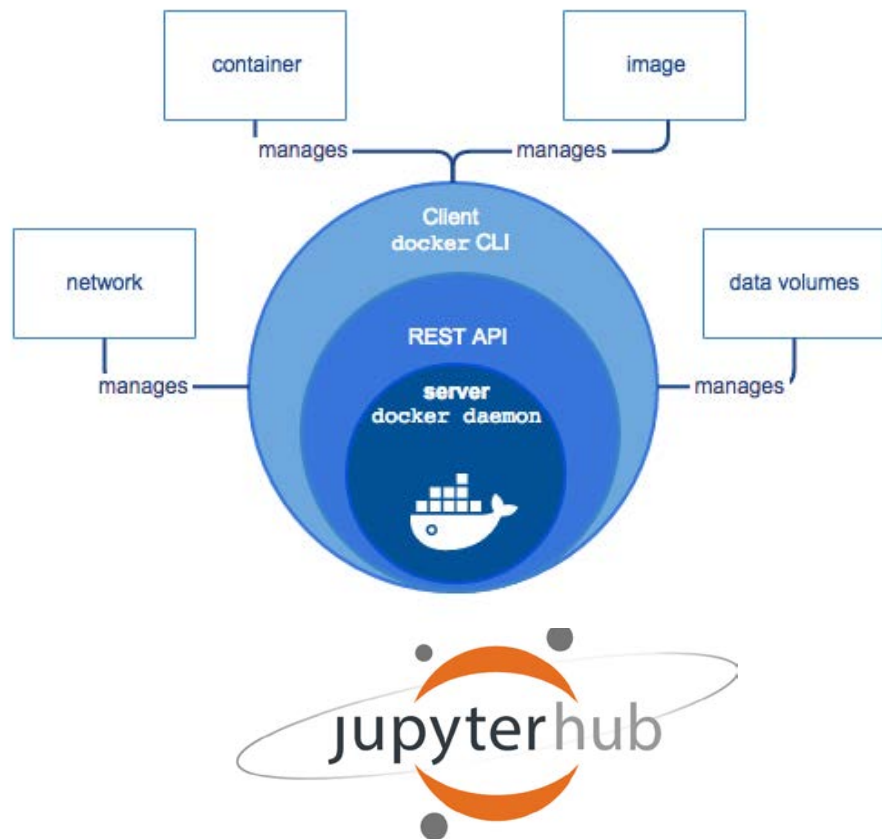
Databases

GUI-query


Notebooks

RAC

Empowered by the open-source ecosystem

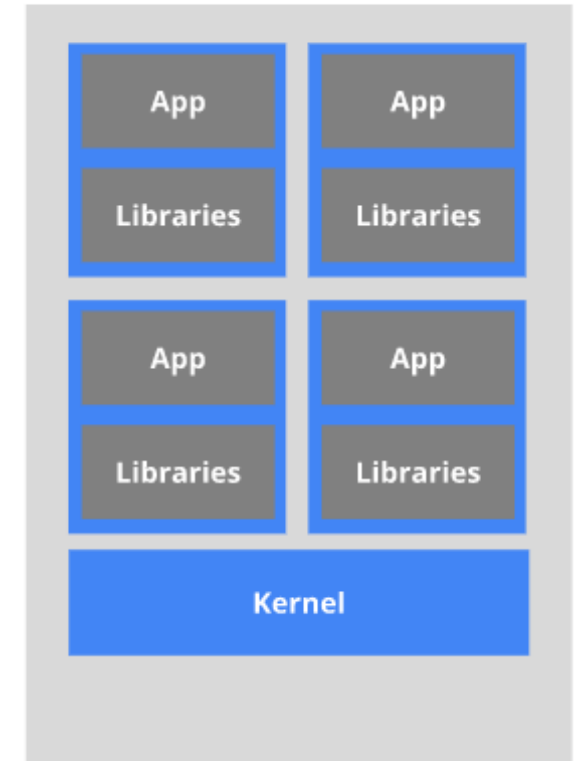


The old way: Applications on host



*Heavyweight, non-portable
Relies on OS package manager*

The new way: Deploy containers



*Small and fast, portable
Uses OS-level virtualization*

Reproducible notebooks on Kubernetes



Turn a Git repo into a collection of interactive notebooks

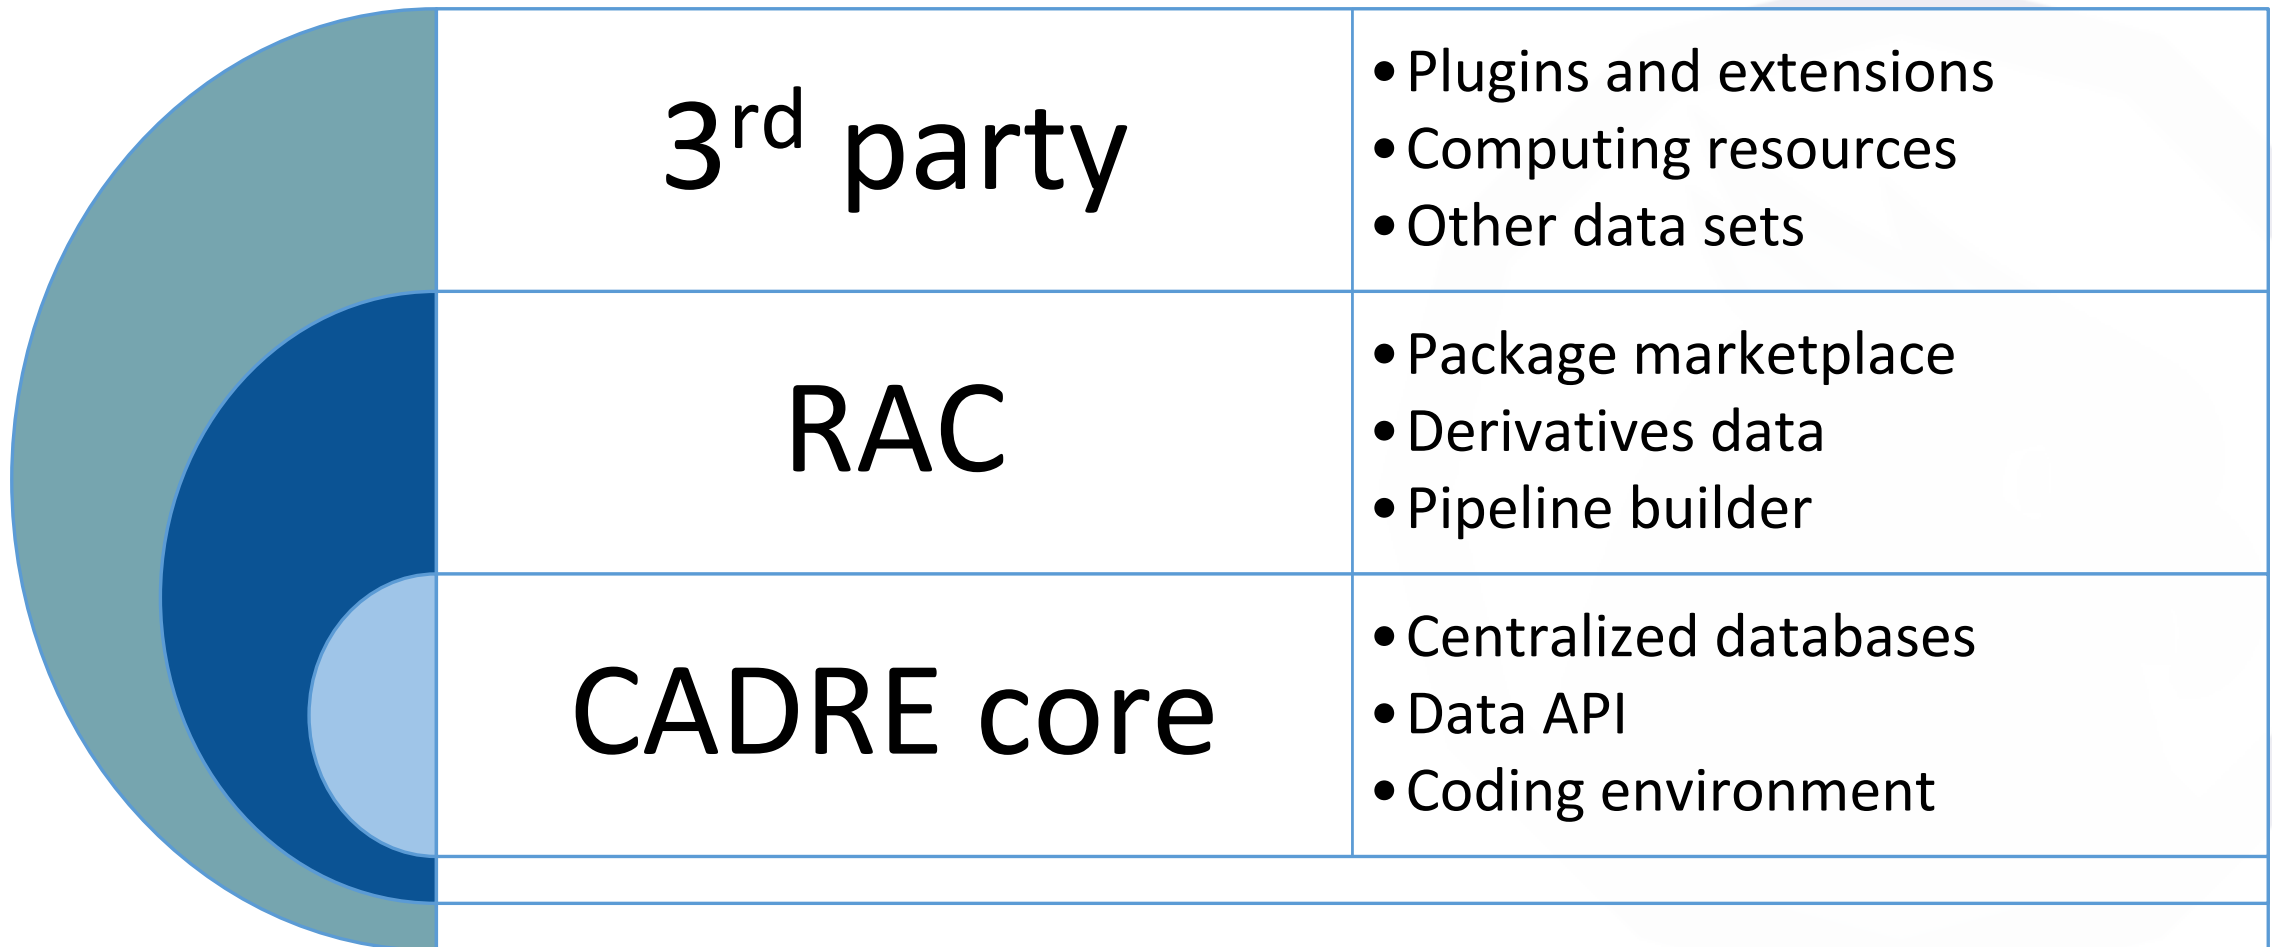
Have a repository full of Jupyter notebooks? With Binder, open those notebooks in an executable environment, making your code immediately reproducible by anyone, anywhere.

Demo 5

<https://github.com/iuni-cadre/ISSI-tutorial>



The CADRE ecosystem



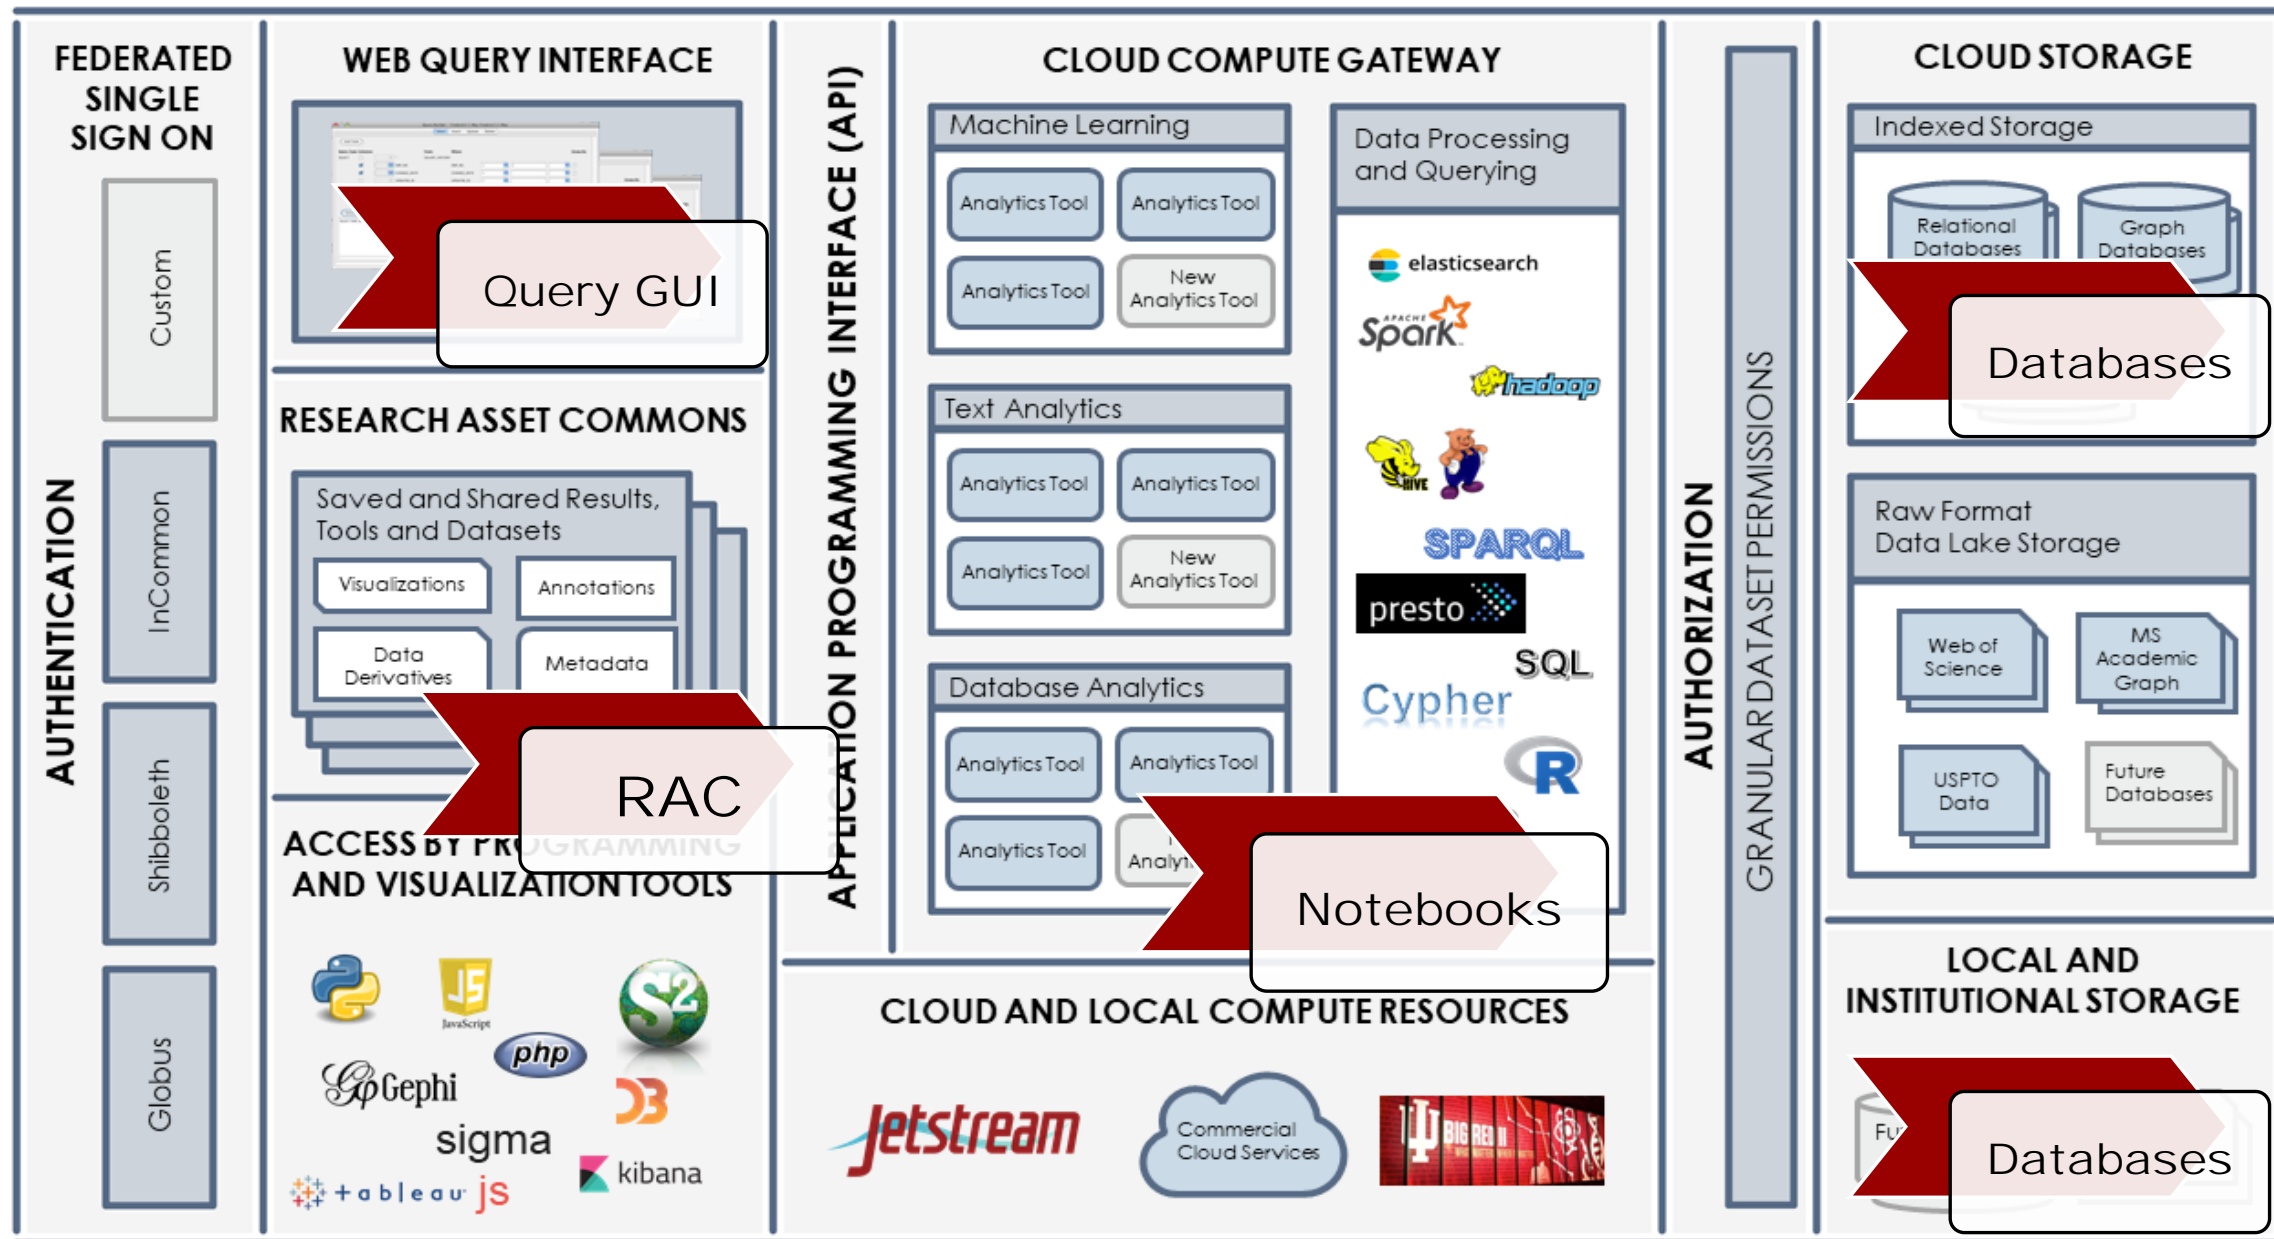
Reproducible notebooks on Kubernetes



Turn a Git repo into a collection of interactive notebooks

Have a repository full of Jupyter notebooks? With Binder, open those notebooks in an executable environment, making your code immediately reproducible by anyone, anywhere.

SHARED BIGDATA-GATEWAY FOR RESEARCH LIBRARIES (SBD-GATEWAY)





Microsoft®
Research



Q&A

The CADRE TEAM




CADRE related events

Apr. 2019



- 2019 CADRE meeting
- CADRE Fellowship open
- 1st Fellows announced
- ISSI workshop & tutorial



Sep. 2019



May. 2020



- 2020 CADRE meeting
- BTAA Library Conference 2020
- 2020 CADRE hackathon



Contact Us



<https://cadre.iu.edu>



cadre@iu.edu



[@CADRE_Project](https://twitter.com/CADRE_Project)